

Contents

The Cozy Packages

1

The Cozy Packages

The standard packages that ship with Cozy: probability distributions, polynomials, finance, a solar almanac, structured random matrices, physical constants, scatter plots, symbolic differentiation, and a guided tour — all written in Cozy itself.

A package is a file of `let` definitions; `load("packages/name.cz")` runs it in the current session and its bindings persist. Records of closures act as namespaces (`norm.cdf`), and packages validate their inputs with `assert`, so they fail like builtins do — with a message, not an accident. Every transcript below is machine-verified against the interpreter by `tests/run_manual.sh`, and every numerical claim is cross-checked in the package's own golden tests against an independent reference (SciPy, NumPy, Python's `datetime`, or the `astral` library — named in each section).

1. `dist.cz` — probability distributions

Six distributions — `norm`, `student`, `chi2`, `fdist`, `expo`, `unif` — each a record with `pdf`, `cdf`, `inv` (quantile), and `rand`. Parameters are explicit: `norm.cdf(x, mu, sigma)`. Quantiles with no closed form use bracket expansion plus `fzero` on the CDF; the quantile domain is $0 < p < 1$. Cross-checked against SciPy (`tests/34_dist.test`).

```
cozy> load("packages/dist.cz")
cozy> norm.cdf(1.96, 0, 1)
0.975002
cozy> norm.inv(0.975, 0, 1)
1.95996
cozy> student.inv(0.975, 30)
2.04227
cozy> chi2.inv(0.95, 3)
7.81473
cozy> fdist.cdf(3.0, 4, 20)
0.956799
```

A complete two-sample t-test — simulate, test, decide:

```
cozy> load("packages/dist.cz")
cozy> rng(42)
cozy> let x1 = norm.rand(40, 100, 15)
[75.8016; 123.017; 111.725; 93.9971; 100.238; 98.0904; 107.158; 90.1486; 90.4082; 94.4609; 96.6851; 112.
cozy> let x2 = norm.rand(40, 108, 15)
[102.897; 92.0654; 112.122; 98.9768; 112.891; 112.931; 92.8931; 102.642; 111.903; 100.646; 100.94; 123.6
cozy> let t = (mean(x2) - mean(x1)) / sqrt(var(x1)/40 + var(x2)/40)
2.46819
cozy> 2 * (1 - student.cdf(abs(t), 78))
0.0157683
```

Out-of-domain probabilities refuse with the package's own message:

```
cozy> load("packages/dist.cz")
error: quantile: p must be in (0,1), got 1.5
cozy> student.inv(1.5, 10)
```

Function	Worked example	Result
<code>norm.pdf/cdf/inv(x, mu, sigma)</code>	<code>norm.cdf(1.96, 0, 1)</code>	0.975002
<code>norm.rand(n, mu, sigma)</code>	<code>mean(norm.rand(500, 10, 2)) > 9</code>	true
<code>student.pdf/cdf/inv(x, v)</code>	<code>student.inv(0.975, 30)</code>	2.04227
<code>chi2.pdf/cdf/inv(x, k)</code>	<code>chi2.inv(0.95, 3)</code>	7.81473
<code>fdist.pdf/cdf/inv(x, d1, d2)</code>	<code>fdist.cdf(3.0, 4, 20)</code>	0.956799
<code>expo.pdf/cdf/inv(x, rate)</code>	<code>expo.inv(0.5, 2)</code>	0.346574
<code>unif.pdf/cdf/inv(x, a, b)</code>	<code>unif.cdf(3, 0, 10)</code>	0.300000

2. poly.cz — polynomials

Coefficients are row vectors, highest power first: `[2, -3, 1]` is $2x^2 - 3x + 1$. Roots come from the companion matrix and `eig` — the same algorithm Octave uses, on Cozy's LAPACK-verified eigensolver. Cross-checked against NumPy (`tests/37_poly.test`).

```
cozy> load("packages/poly.cz")
cozy> polyval([2, -3, 1], 4)
21
cozy> roots([1, -6, 11, -6])'
[1, 2, 3]
cozy> roots([1, 0, 1])'
[1i, -1i]
cozy> companion([1, -6, 11, -6])
[6, -11, 6; 1, 0, 0; 0, 1, 0]
```

Least-squares fitting (Vandermonde + backslash), and calculus that inverts:

```
cozy> load("packages/poly.cz")
cozy> let x = [0, 1, 2, 3, 4]
[0, 1, 2, 3, 4]
cozy> let y = [1.1, 2.9, 9.2, 19.1, 32.8]
[1.1, 2.9, 9.2, 19.1, 32.8]
cozy> polyfit(x, y, 2)
[1.95714, 0.131429, 1.01429]
cozy> polyder([1, -6, 11, -6])
[3, -12, 11]
cozy> polyint(polyder([1, -6, 11, -6]), -6)
[1, -6, 11, -6]
cozy> conv([1, -1], [1, -2])
[1, -3, 2]
```

Function	Worked example	Result
<code>polyval(c, x)</code>	<code>polyval([2, -3, 1], 4)</code>	21
<code>roots(c)</code>	<code>sort(real(roots([1, -6, 11, -6])))'</code>	[1.00000, 2.00000, 3.00000]
<code>companion(c)</code>	<code>companion([1, 0, -4])</code>	[0.00000, 4.00000; 1.00000, 0.00000]
<code>polyfit(x, y, n)</code>	<code>polyfit([1, 2, 3], [3, 7, 13], 2)</code>	[1.00000, 1.00000, 1.00000]
<code>polyder(c)</code>	<code>polyder([1, -6, 11, -6])</code>	[3, -12, 11]

Function	Worked example	Result
<code>polyint(c, k)</code>	<code>polyint([3, -12, 11], -6)</code>	<code>[1.00000, -6.00000, 11.0000, -6.00000]</code>
<code>conv(a, b)</code>	<code>conv([1, -1], [1, -2])</code>	<code>[1.00000, -3.00000, 2.00000]</code>

3. `finance.cz` — the HP-12C's greatest hits

Time value of money, cash flows, bonds, amortization, and date arithmetic. Sign convention (HP-12C): money received is positive, money paid is negative; rates are per period, as decimals. Cross-checked against SciPy and Python's `datetime` (`tests/38_finance.test`).

The mortgage block — payment, implied rate, savings growth:

```
cozy> load("packages/finance.cz")
cozy> pmt(360, 0.005, 250000, 0)
-1498.88
cozy> rate(360, 250000, -1498.876313, 0)
0.005
cozy> fv(120, 0.004, 0, -200)
30726.4
```

Cash flows and bonds:

```
cozy> load("packages/finance.cz")
cozy> npv(0.10, [-1000, 300, 420, 680])
130.729
cozy> irr([-1000, 300, 420, 680])
0.163406
cozy> bond_price(100, 0.08, 0.06, 10, 1)
114.72
cozy> bond_ytm(114.7202, 100, 0.08, 10, 1)
0.06
cozy> bond_mduration(100, 0.08, 0.06, 10, 1)
7.0236
```

`amort` returns the whole schedule as a record of columns — a small data frame you can index, sum, and plot:

```
cozy> load("packages/finance.cz")
cozy> let s = amort(1000, 0.01, 3)
{period = [1; 2; 3], payment = [340.022; 340.022; 340.022], interest = [10; 6.69978; 3.36656], principal
cozy> fields(s)'
["period", "payment", "interest", "principal", "balance"]
cozy> s.payment[1]
340.022
cozy> s.balance'
[669.978, 336.656, 4.26326e-12]
cozy> sum(s.principal)
1000
```

Dates, HP-12C style — actual day counts, date arithmetic, day of week, and the bond market's 30/360 convention. `datestr` returns a display string, `daterec` the `{y, m, d}` record, and `today()` a serial day number, so arithmetic reads naturally — `datestr(today() + 90)` is the date in 90 days:

```
cozy> load("packages/finance.cz")
cozy> days(2026, 1, 1, 2026, 7, 9)
189
```

```

cozy> datestr(datenum(2026, 7, 17) + 90)
"2026-10-15"
cozy> daterec(datenum(2024, 2, 29))
{y = 2024, m = 2, d = 29}
cozy> dateadd(2024, 2, 28, 2)
{y = 2024, m = 3, d = 1}
cozy> dow(2026, 7, 9)
4
cozy> days360(2024, 1, 31, 2024, 7, 31)
180
cozy> daterec(today()).y >= 2026
true

```

Function	Worked example	Result
pmt(n, i, pv, fv)	pmt(360, 0.005, 250000, 0)	-1498.88
pv(n, i, pmt, fv)	pv(360, 0.005, -1498.88, 0)	250001.
fv(n, i, pv, pmt)	fv(120, 0.004, 0, -200)	30726.4
nper(i, pv, pmt, fv)	nper(0.005, 250000, -1498.88, 0)	359.998
rate(n, pv, pmt, fv)	rate(360, 250000, -1498.876313, 0)	0.00500000
npv(r, cfs)	npv(0.10, [-1000, 300, 420, 680])	130.729
irr(cfs)	irr([-1000, 300, 420, 680])	0.163406
bond_price(face, crate, y, n, freq)	bond_price(100, 0.08, 0.06, 10, 1)	114.720
bond_ytm(price, face, crate, n, freq)	bond_ytm(114.7202, 100, 0.08, 10, 1)	0.0600000
bond_duration(face, crate, y, n, freq)	bond_duration(100, 0.08, 0.06, 10, 1)	7.44502
bond_mduration(face, crate, y, n, freq)	bond_mduration(100, 0.08, 0.06, 10, 1)	7.02360
bond_convexity(face, crate, y, n, freq)	bond_convexity(100, 0.08, 0.06, 10, 1)	65.1716
amort(principal, i, n)	amort(1000, 0.01, 3).balance[3]	4.26326e-12
datenum(y, m, d)	datenum(2026, 7, 17)	2.46124e+06
datestr(jdn)	datestr(datenum(2026, 7, 17) + 90)	"2026-10-15"
daterec(jdn)	daterec(datenum(2026, 7, 17))	{y = 2026.00, m = 7.00000, d = 17.0000}
dateadd(y, m, d, k)	dateadd(2024, 12, 31, 1)	{y = 2025.00, m = 1.00000, d = 1.00000}
today()	today() == datenum(now().y, now().m, now().d)	true
dow(y, m, d)	dow(2026, 7, 17)	5.00000
days(y1,m1,d1, y2,m2,d2)	days(2026, 1, 1, 2026, 7, 17)	197.000
days360(y1,m1,d1, y2,m2,d2)	days360(2024, 1, 31, 2024, 7, 31)	180

4. astro.cz — the solar almanac

Sunrise, sunset, three grades of twilight, solar noon, day length, sun position, and moon phase, for any latitude and longitude (degrees; north and east positive). Times are local decimal hours given a UTC offset `tz`; `hm` renders them as “HH:MM”. NOAA solar position algorithm, within about a minute of the `astral` reference library (`tests/39_astro.test`). A `places` record preloads coordinates for several cities.

```
cozy> load("packages/astro.cz")
cozy> hm(sunrise(places.alexandria_va.lat, places.alexandria_va.lon, 2026, 7, 10, -4))
"05:52"
cozy> hm(sunset(places.duluth_ga.lat, places.duluth_ga.lon, 2026, 7, 10, -4))
"20:50"
cozy> hm(solar_noon(places.ljubljana.lon, 2026, 7, 10, 2))
"13:07"
cozy> day_length(places.ljubljana.lat, places.ljubljana.lon, 2026, 7, 10)
15.5331
```

Planning a drive to arrive before dark — civil dawn at the origin to civil dusk at the destination:

```
cozy> load("packages/astro.cz")
cozy> let w = drive_daylight(places.alexandria_va, places.duluth_ga, 2026, 7, 10, -4)
{depart_dawn = 5.34382, depart_rise = 5.86728, arrive_set = 20.8376, arrive_dusk = 21.3161, window_hours}
cozy> hm(w.depart_dawn) + " to " + hm(w.arrive_dusk)
"05:21 to 21:19"
cozy> w.window_hours
15.9723
```

Sun position, moon illumination, and the honest polar refusal:

```
cozy> load("packages/astro.cz")
error: the sun does not cross 90.833 degrees at latitude 80 on 2026-6-21
cozy> let p = sun_position(38.8048, -77.0469, 2026, 7, 10, 13.22, -4)
{alt = 73.3594, az = 179.65}
cozy> p.alt
73.3594
cozy> moon_illum(2026, 7, 10)
0.19449
cozy> sunrise(80, 0, 2026, 6, 21, 0)
```

Function	Worked example	Result
<code>sunrise(lat, lon, y, m, d, tz)</code>	<code>hm(sunrise(38.8048, -77.0469, 2026, 7, 17, -4))</code>	"05:57"
<code>sunset(lat, lon, y, m, d, tz)</code>	<code>hm(sunset(38.8048, -77.0469, 2026, 7, 17, -4))</code>	"20:31"
<code>dawn_civil / dusk_civil(lat, lon, y, m, d, tz)</code>	<code>hm(dawn_civil(38.8048, -77.0469, 2026, 7, 17, -4))</code>	"05:26"
<code>dawn_nautical / dusk_nautical(lat, lon, y, m, d, tz)</code>	<code>hm(dawn_nautical(38.8048, -77.0469, 2026, 7, 17, -4))</code>	"04:48"
<code>dawn_astro / dusk_astro(lat, lon, y, m, d, tz)</code>	<code>hm(dusk_astro(38.8048, -77.0469, 2026, 7, 17, -4))</code>	"22:23"
<code>solar_noon(lon, y, m, d, tz)</code>	<code>hm(solar_noon(-77.0469, 2026, 7, 17, -4))</code>	"13:14"
<code>day_length(lat, lon, y, m, d)</code>	<code>day_length(46.0569, 14.5058, 2026, 7, 17)</code>	15.3443

Function	Worked example	Result
<code>sun_position(lat, lon, y, m, d, hour, tz)</code>	<code>sun_position(38.8048, -77.0469, 2026, 7, 17, 13.2, -4).alt</code>	72.2914
<code>moon_age(y, m, d)</code>	<code>moon_age(2026, 7, 17)</code>	2.70686
<code>moon_illum(y, m, d)</code>	<code>moon_illum(2026, 7, 17)</code>	0.0806581
<code>hm(h)</code>	<code>hm(6.05)</code>	"06:03"
<code>places</code>	<code>places.duluth_ga.lat</code>	34.0029
<code>drive_daylight(from, to, y, m, d, tz)</code>	<code>drive_daylight(places.alexandria5v8284, places.duluth_ga, 2026, 7, 17, -4).window_hours</code>	15.284

5. `rmt.cz` — random matrices, structured

`packages/rmt.cz` is sugar over `randn`/`qr`/`eye` for the matrices you actually reach for at the prompt: symmetric, positive definite (chol-safe by construction: eigenvalues at least 1), Haar orthogonal, permutations and their matrices, correlation, row-stochastic (random Markov chains), and the Gaussian orthogonal ensemble — whose spectrum follows Wigner's semicircle on $[-2, 2]$, as the last line demonstrates on a 200×200 draw. Everything is reproducible under `rng(seed)`.

```
cozy> load("packages/rmt.cz")
cozy> format(6)
cozy> rng(11)
cozy> randorth(2)
[0.684752, 0.728776; -0.728776, 0.684752]
cozy> let P = randspd(3); chol(P)[1, 1] > 0
true
cozy> randperm(6)
[4, 5, 1, 2, 3, 6]
cozy> permmat([3, 1, 2])
[0, 0, 1; 1, 0, 0; 0, 1, 0]
cozy> let H = goe(200); max(abs(eig(H).values))
1.99476
```

Function	Worked example	Result
<code>randsym(n)</code>	<code>let S = randsym(3); max(max(abs(S - S')))) == 0</code>	true
<code>randspd(n)</code>	<code>let P = randspd(3); min(eig(P).values) >= 1</code>	true
<code>wishart(n)</code>	<code>let W = wishart(3); min(eig(W).values) > 0</code>	true
<code>randorth(n)</code>	<code>let Q = randorth(3); max(max(abs(Q' * Q - eye(3)))) < 1e-12</code>	true
<code>randperm(n)</code>	<code>sort(randperm(5)) == 1:5</code>	[true, true, true, true, true]
<code>permmat(p)</code>	<code>permmat([2, 3, 1])</code>	[0, 1, 0; 0, 0, 1; 1, 0, 0]
<code>randcorr(n)</code>	<code>let C = randcorr(3); max(abs(diag(C) - 1)) < 1e-12</code>	true

Function	Worked example	Result
<code>randstoch(n)</code>	<code>let T = randstoch(4); max(abs(sum(T, 2) - 1)) < 1e-12</code>	<code>true</code>
<code>goe(n)</code>	<code>let H = goe(150); max(abs(eig(H).values)) < 2.4</code>	<code>true</code>

6. `phys.cz` — physical constants

`packages/phys.cz` is the CODATA 2018 constants as a record — exact where the 2019 SI redefinition makes them exact (`c`, `h`, `k`, `NA`, `qe`), the recommended values elsewhere. The transcript computes Earth's surface gravity from `G`, recovers the speed of light from Maxwell's relation `1/sqrt(eps0*mu0)`, and expresses thermal energy at room temperature in electron volts:

```
cozy> load("packages/phys.cz")
cozy> format(6)
cozy> phys.c
299792458
cozy> phys.hbar
1.05457e-34
cozy> phys.G * 5.972e24 / 6.371e6 ^ 2
9.81997
cozy> 1 / sqrt(phys.eps0 * phys.mu0)
2.99792e+08
cozy> phys.k * 300 / phys.eV
0.0258520
```

Function	Worked example	Result
<code>phys.c</code>	<code>phys.c == 299792458</code>	<code>true</code>
<code>phys.h</code>	<code>phys.h</code>	<code>6.62607e-34</code>
<code>phys.hbar</code>	<code>abs(phys.hbar * 2 * pi - phys.h) < 1e-45</code>	<code>true</code>
<code>phys.G</code>	<code>phys.G</code>	<code>6.67430e-11</code>
<code>phys.k (thermal, in eV)</code>	<code>phys.k * 300 / phys.eV</code>	<code>0.0258520</code>
<code>phys.NA</code>	<code>phys.NA == 6.02214076e23</code>	<code>true</code>
<code>phys.R</code>	<code>phys.R</code>	<code>8.31446</code>
<code>phys.qe</code>	<code>phys.qe</code>	<code>1.60218e-19</code>
<code>phys.alpha</code>	<code>1 / phys.alpha</code>	<code>137.036</code>
<code>phys.sigma</code>	<code>phys.sigma</code>	<code>5.67037e-08</code>
<code>phys.ly / phys.au</code>	<code>phys.ly / phys.au</code>	<code>63241.1</code>

7. `scatter.cz` — scatter plots without touching the core

Proof that the frozen language plots more than it appears to: `plot()`'s `style = "points"` option is honored by every backend (SVG circles in the browser, point markers in `ascii` and `gnuplot`), so scatter plots are a pure package. `scatter(x, y)` and `scatter_titled(x, y, t)` wrap that fact; `jitter(x, amount)` spreads overplotted values for legibility. Per-point sizes and colors would need backend changes and are deliberately absent — the package's header says so. The transcript exercises the numeric helper; the SVG output itself is asserted by the plot test suite:

```
cozy> load("packages/scatter.cz")
cozy> rng(1); let j = jitter(1:100, 0.1);
cozy> max(abs(j - (1:100))) <= 0.05
```

```
true
cozy> size(jitter(rand(3, 4), 0.2))
[3, 4]
```

Function	Worked example	Result
jitter(x, a) stays within a/2	max(abs(jitter(1:100, 0.1) - (1:100))) <= 0.05	true
jitter preserves shape	size(jitter(rand(3, 4), 0.2)) == [3, 4]	[true, true]
mean displacement is small	abs(mean(jitter(zeros(1, 2000), 0.2))) < 0.01	true

8. symb.cz — symbolic differentiation

A symbolic differentiator in pure Cozy, addressed in mathematics and answering in it. The front door is a string:

```
cozy> format(6)
cozy> load("packages/symb.cz")
cozy> deriv("sin(x)/x")
"((cos(x) / x) - (sin(x) / x^2))"
cozy> deriv("x^x")
"(exp((x * log(x))) * (log(x) + (x / x)))"
cozy> deriv("exp(-x^2)")
"(-2 * (exp((-x^2)) * x))"
cozy> let f = dfun("sin(x)/x");
cozy> fzero(f, 4, 5)
4.49341
cozy> fminbnd(ffun("sin(x)/x"), 4, 5)
{x = 4.49341, fx = -0.217234}
cozy> integral(dfun("x^3"), 1, 2)
7.00000
```

`deriv` parses, differentiates, simplifies, and prints; `ffun` and `dfun` lift a string to an ordinary function (of the original, or of its derivative), which then rides the whole numeric ecosystem — the transcript above finds the $\tan x = x$ critical point of $\sin(x)/x$ twice over (`fzero` on the symbolic derivative, `fminbnd` on the original, agreeing at 4.49341) and checks the Fundamental Theorem of Calculus (the integral of `dfun("x^3")` over $[1, 2]$ is 7, the function's change).

How it works. Expressions are nested records; `ddx` differentiates by structural recursion. `sub` and `divx` desugar into `add/mul/negative` powers at construction, so the product, power, and chain rules alone carry the whole calculus — the quotient rule falls out of `d(b-1)` for free, which is why `deriv("sin(x)/x")` above prints the textbook form. The parser is recursive descent over `s[i]` (character classes are chained comparisons: `"0" <= c <= "9"`), unary minus binds looser than `^` so `-x^2` means $-(x^2)$, and general powers desugar as `f^g = exp(g · log f)` — which is how `x^x` differentiates with no special case. The printer recognizes division and subtraction; `simp` folds and hoists constants. Since v2.19.2 the dispatchers are written with `elseif` — the package whose nine-deep end-cascades helped motivate the feature was first to enjoy it.

The constructor layer remains public — it *is* the representation, and remains the door for programmatic tree-building: `C(5)`, `X`, `add`, `mul`, `powc`, `sinx`, ... plus `evalx` (evaluate at a point), `subst` (compose), `dn` (*n*-th derivative), and `taylor` (series coefficients by repeated `ddx`). The transcript cross-checks the symbolic answer against numeric differencing — two independent derivatives agreeing inside one verified session — and recovers the sine series $(0, 1, 0, -1/6, \dots)$ from pure record recursion:

```
cozy> format(4)
cozy> load("packages/symb.cz")
```

```

cozy> let e = add(powc(X, 3), mul(C(5), sinx(X)));
cozy> show(simp(ddx(e)))
"((3 * x^2) + (5 * cos(x)))"
cozy> let d = fn f -> fn x -> (f(x + h) - f(x - h)) / (2 * h) where h = 1e-6
<fn/1>
cozy> abs(evalx(ddx(e), 2) - d(fn t -> t ^ 3 + 5 * sin(t))(2)) < 1e-8
true
cozy> show(dn(powc(X, 5), 3))
"(60 * x^2)"
cozy> taylor(sinx(X), 7)
[0.000, 1.000, -0.000, -0.1667, 0.000, 0.008333, -0.000, -0.0001984]

```

(History, kept honest: v2.12.0 shipped constructor-only under a recorded limitation — “no substring access” — that v2.13.1 discovered was false all along: strings index like arrays, golden-tested since the strings phase. v2.14.0 built the parser that was always possible; v2.15.0 taught the printer textbook manners; v2.16.0 added the function-space lifts. The wrong entry and its correction are both in `KNOWN_LIMITATIONS` — the goldens outrank the maintainer’s memory.)

Function	Worked example	Result
d/dx of x ³ at 2 is 12	evalx(ddx(powc(X, 3)), 2) == 12	true
chain rule through exp(2x)	abs(evalx(ddx(subst(expx(X), mul(C(2), X))), 0) - 2) < 1e-12	true
taylor of exp begins 1, 1, 1/2	max(abs(taylor(expx(X), 2) - [1, 1, 0.5])) < 1e-9	true
parse then evaluate	abs(evalx(parse("2*pi"), 0) - 2 * pi) < 1e-12	true
deriv: string in, string out	deriv("x^2") == "(2 * x)"	true
dfun: derivative as a function	abs(dfun("x^3")(2) - 12) < 1e-12	true
FTC: integral of dfun is the change	abs(integral(dfun("x^3"), 1, 2) - 7) < 1e-6	true

9. demo.cz — the tour

Not a library: a performance. `load("packages/demo.cz")` plays a guided tour in six acts — executable mathematics (Basel, the Gaussian integral), the blackboard’s word order (**where**, chained masks), functions as values (anonymous application, Euclid’s recursion, n-fold self-composition converging on the golden ratio), the pipelines act with both crown jewels (the oscillation pipe — values flavor-changing between `~>` and `|>` down one chain — and the fan-out dashboard), the calculus (`symb.cz`’s `deriv` and a symbolic derivative fed to `fzero`), money (the HP-12C mortgage), and a plotted finale. Every demonstration is a `SINGLE SOURCE STRING` — printed as the caption and executed by `eval` — so what the tour shows is, byte for byte, what it runs; drift between shown and real code is impossible by construction. The pacing is deliberate: `pause()` waits for Enter after each feature at a live prompt — including the browser, where the pending output is streamed to the terminal first and the interpreter then truly waits for Enter (`Asyncify`) — while piped input or `make test` drains the pauses instantly at EOF. One file, every temperament. Seeded where random, deterministic throughout. Point a newcomer at this file first.

```
cozy> load("packages/demo.cz")
```

10. sparselin.cz — iterative sparse linear algebra

The design’s answer to `S \ b` and `eig(S)`: solvers as packages on the founding kernel `S * v`, in pure Cozy. `cg(A, b)` solves symmetric positive definite systems by conjugate gradient (returns `{x, iters, relres}`);

zero start, relative tolerance 1e-10); `powerit(A)` finds the dominant eigenpair `{value, vector, iters}` by power iteration with a Rayleigh-quotient estimate (deterministic start; `powerit_from(A, x0)` to choose your own). Both also accept dense matrices — anything with `*` and a column vector works. The dense-only kernels gate with pointers here: `eig(S)` and friends name `sparselin.cz` in their error messages.

```
cozy> load("packages/sparselin.cz")
cozy> let A = sparse([4, 1, 0; 1, 3, 1; 0, 1, 5]); let s = cg(A, [1; 2; 3]); s.x
[0.137255; 0.45098; 0.509804]
cozy> s.iters
3
cozy> powerit(sparse([2, 1; 1, 2])).value
3
cozy> rng(7); let W = sprandn(200, 200, 0.02); let S = W + W' + 10 * speye(200); cg(S, ones(200, 1)).rel
true
```

CG's finite-termination property is visible above: a 3x3 system converges in exactly 3 iterations. The 200x200 line is the trigger workload the sparse design waited for, solved without ever forming a dense matrix.

11. autodiff.cz — forward-mode automatic differentiation

Derivatives as exact language values (design entry 4a). `d(f)` returns the derivative of a scalar function as a function; `grad(f)` returns the gradient of $f : \mathbb{R}^n \rightarrow \mathbb{R}$, one dual pass per component. Both are machine-exact — dual numbers apply the chain rule through the numeric tower itself, so closures, conditionals, and composition all just work. This is the gradient infrastructure the optimization capability builds on.

```
cozy> load("packages/autodiff.cz")
cozy> let df = d(fn x -> x^3 - 2*x); df(2)
10
cozy> d(abs)(-3)
-1
cozy> grad(fn x -> sum(x .* x))([1.0; 2.0; 3.0])
[2; 4; 6]
cozy> let rosen = fn x -> (1 - x[1])^2 + 100 * (x[2] - x[1]^2)^2; grad(rosen)([1.0; 1.0])
[0; 0]
cozy> grad(rosen)([0.0; 0.0])
[-2; 0]
```

`hess(f)` returns the exact Hessian, one hyper-dual pass per index pair with symmetry filled in:

```
cozy> load("packages/autodiff.cz")
cozy> hess(fn x -> x[1] * x[2]^2)([2.0; 3.0])
[0, 6; 6, 4]
```

Discussion. The Rosenbrock gradient vanishing at `[1; 1]` is the textbook minimum, found here without a single hand-written derivative. `grad` seeds one coordinate direction per pass with `dual(x, e_i)`, and `dual eps` of the result is the exact partial — the `d` one-liner is `fn f -> fn x -> dual eps(f(dual(x, 1)))`, the whole engine in one line.

12. optim.cz — multivariate optimization

Design entry 3, on autodiff's exact gradients: BFGS with Armijo backtracking. `minimize(f, x0)` and `maximize(f, x0)` (maximization is negation) return `{x, fx, iters, converged}`. `minimize_box / maximize_box(f, x0, lb, ub)` handle bounds by gradient projection. `minimize_con / maximize_con(f, x0, cons)` handle general constraints by augmented Lagrangian — `cons` is a record with `eq` (driven to 0) and/or `ineq` (driven ≤ 0), each a function returning a column — and the inner solver is `minimize` itself. Objectives must flow dual numbers: elementwise ops, `*`, and reductions (see the manual's dual section).

```
cozy> load("packages/optim.cz")
```

```

cozy> minimize(fn x -> (1 - x[1])^2 + 100 * (x[2] - x[1]^2)^2, [-1.2; 1.0]).x
[1; 1]
cozy> let cd = fn x -> x[1]^0.3 * x[2]^0.7
cozy> let bud = fn x -> [2 * x[1] + 3 * x[2] - 12]
cozy> maximize_con(cd, [1.0; 1.0], {eq = bud}).x
[1.8; 2.8]

```

`minimize_newton` / `maximize_newton(f, x0)` run true Newton on the exact gradient and Hessian, with Levenberg damping when the Hessian is indefinite. On a positive-definite quadratic the first step lands on the minimum:

```

cozy> load("packages/optim.cz")
cozy> minimize_newton(fn x -> sum(x .* ([3.0, 1.0; 1.0, 2.0] * x)) - sum([1.0; 1.0] .* x), [9.0; -7.0]).
1

```

Discussion. The last transcript is a Cobb-Douglas utility maximum on a budget line — the constrained maximization the design was scoped around — and `[1.8; 2.8]` is the analytic answer (`a m / p`, `b m / q`) to the digits shown, with no hand-written derivative or Lagrangian algebra anywhere: `grad` differentiates the augmented objective, `kink` and `all`, because `max` on duals takes the one-sided derivative.

Writing your own

The pattern the twelve packages follow: a file of `let` definitions, `assert`-validated inputs, and a golden test file whose reference values come from an independent implementation. Multi-line definitions are fine wherever a bracket is open. `save` your workspace, `load` it tomorrow — a package is just a workspace you chose to keep.

The namespace law (design entry 7 — records are the module system). A package that wants a namespace packs its public API into a record: `{cdf = fn ..., pdf = fn ...}` — `fields(pkg)` is then the manifest and `getfield(pkg, name)` the dynamic door, and sibling fields may call each other through the record’s own global name (functions resolve globals at CALL time, so `p.isod` inside `p.isev` works). The same late binding is also the one trap: **a record namespace hides the face, never the body**. Helpers that your public functions call remain ordinary globals forever — pack the faces into a record, `keep()` the record alone, and every call dies with “undefined name”. So: prefix internal helpers with the package’s tag (`op_`, `sl_`, `ad_`), document that `keep()` must spare them, and never expect the record to encapsulate. Small instrument packages with a handful of daily-typed faces (`minimize`, `d`, `grad`, `cg`) stay flat by deliberate choice — brevity is the feature — with prefixed helpers; a large API is where the record namespace earns its keep.